

## KNN Algorithm

The k-nearest neighbor algorithm, sometimes referred to as KNN or k-NN, is a nonparametric supervised learning classifier that relies on closeness to assign labels to or predict the grouping of a single data point.

Although it can be applied to regression or classification issues, it is typically employed as a classification technique because it is predicated on the idea that similar points can be identified nearby.

A majority vote is used to apply a class label to classification problems, meaning the label that is most frequently expressed around a particular data point is utilized. Although this counts as "majority voting," the term "majority vote" is more frequently used in literature. These terms differ in that a "majority vote" actually requires a majority of more than 50%, which only applies when there are only two options. When there are many classes, such as four categories, deciding on a class does not always require a majority of votes; a class label can be assigned with a vote of more than 25%.

Classification difficulties use a notion similar to regression problems, except in this situation, a classification prediction is made using the average of the k nearest neighbors. Classification is utilized for discrete data whereas regression is used for continuous values in this situation. However, defining the distance is necessary before a categorization can be determined. The most popular measurement is euclidean distance, and we shall go into more detail later.

Noting that the KNN algorithm just saves a training data set rather than going through a training phase, it should be noted that it belongs to a family of "lazy learning" models. This also implies that whenever a classification or prediction is made, all calculation takes place. It is also known as an instance-based or memory-based learning approach since it stores all of its training data entirely in memory.

## General Pseudocode

```
Function kNN_Classify(dataset: Array of Example,
    unlabeled: Example,
    k: Integer) -> Label:
    Step 1: Compute distances between the unlabeled example and all examples in the dataset
    distances = []
    For each example in dataset:
        distance = ComputeDistance(unlabeled, example)
        distances.append((distance, example))

    Step 2: Sort the distances in ascending order
    Sort(distances)

    Step 3: Select the k smallest distances
    kNearestNeighbors = distances[0:k]

    Step 4: Get the labels of the k nearest neighbors
    kLabels = []
    For each (_, example) in kNearestNeighbors:
        kLabels.append(example.label)

    Step 5: Determine the most common label among the k neighbors
    predictedLabel = GetMostCommonLabel(kLabels)

    Return predictedLabel
```

```
Function ComputeDistance(a: Example, b: Example) -> Float:
    This can be any distance metric, e.g., Euclidean distance
```

$$d(x, y) = \sqrt{\sum_{i=1}^n (y_i - x_i)^2}$$

```
Return distance
```

```
Function GetMostCommonLabel(labels: Array of Label) -> Label:
    Count the frequency of each label
    labelCounts = CountFrequencies(labels)

    Sort labels by frequency
    sortedLabels = SortByFrequency(labelCounts)

    Return the label with the highest frequency
    Return sortedLabels[0]
```

## ▼ Pseudocode

1. Import necessary libraries

- Data handling and manipulation: pandas, numpy
  - Plotting: matplotlib, seaborn
  - Machine Learning: sklearn
2. Load the dataset
    - Assume that the dataset is loaded into a variable named 'dataframe'
  3. Display general statistics of the dataset
    - Print dataframe.describe()
  4. Prepare the features and target variable
    - `X = dataframe[['wordcount','sentimentValue']]`
    - `y = dataframe['Star Rating']`
  5. Split the dataset into training and testing sets
    - Use train\_test\_split method from sklearn and split X, y
    - Assign the result to X\_train, X\_test, y\_train, y\_test
  6. Scale the features
    - Create a MinMaxScaler object
    - Fit and transform X\_train using the scaler
    - Transform X\_test using the same scaler
  7. Choose the appropriate 'k' value for K-NN
    - Iterate k from 1 to 20
      - Create a KNeighborsClassifier with k neighbors
      - Fit the classifier on the training set
      - Compute the accuracy on the test set and store it
    - Plot the accuracies for all k values
  8. Create a K-NN classifier with the chosen 'k' value, e.g., k=7
    - Fit the classifier on the training set
    - Print the accuracy on both training and test sets
  9. Display confusion matrix and classification report
    - Predict the classes of the test set
    - Print the confusion matrix and classification report
  10. Visualize the decision boundaries
    - Create a mesh grid of the feature space
    - Predict the class for each point in the mesh grid
    - Plot the decision boundaries with different colors for each class
    - Plot the training points on the same plot with corresponding colors
    - Add legends and title to the plot
    - Display the plot

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import matplotlib.patches as mpatches
import seaborn as sb

%matplotlib inline
plt.rcParams['figure.figsize'] = (16, 9)
plt.style.use('ggplot')

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

```
# Here we see the dataset to be worked
dataframe = pd.read_csv(r"reviews_sentiment.csv", sep=';')
dataframe.head(10)
```

	Review Title	Review Text	wordcount	titleSentiment	textSentiment	Star Rating	sentimentValue
0	Sin conexión	Hola desde hace algo más de un mes me pone sin...	23	negative	negative	1	-0.486389
1	faltan cosas	Han mejorado la apariencia pero no	20	negative	negative	1	-0.586187
2	Es muy buena lo recomiendo	Andres e puto amoooo	4	NaN	negative	1	-0.602240
3	Version antigua	Me gustana mas la version anterior esta es mas...	17	NaN	negative	1	-0.616271
4	Esta bien	Sin ser la biblia.... Esta bien	6	negative	negative	1	-0.651784
5	Buena	Nada del otro mundo pero han mejorado mucho	8	positive	negative	1	-0.720443
6	De gran ayuda	Lo malo q necesita de ....pero la app es muy buena	23	positive	negative	1	-0.726825

```
# We se general Statistics
dataframe.describe()
```


	wordcount	Star Rating	sentimentValue
count	257.000000	257.000000	257.000000
mean	11.501946	3.420233	0.383849
std	13.159812	1.409531	0.897987
min	1.000000	1.000000	-2.276469
25%	3.000000	3.000000	-0.108144
50%	7.000000	3.000000	0.264091
75%	16.000000	5.000000	0.808384
max	103.000000	5.000000	3.264579

```
# We define our X and Y, as test as well as train
X = dataframe[['wordcount', 'sentimentValue']].values
y = dataframe['Star Rating'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
#This code help to us to vizualize wich K is better for the problem
k_range = range(1, 20)
scores = []
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors = k)
    knn.fit(X_train, y_train)
    scores.append(knn.score(X_test, y_test))
plt.figure()
plt.xlabel('k')
plt.ylabel('accuracy')
plt.scatter(k_range, scores)
plt.xticks([0,5,10,15,20])
```

```
([<matplotlib.axis.XTick at 0x7ce0dc7eada0>,
<matplotlib.axis.XTick at 0x7ce0dc7eae30>,
<matplotlib.axis.XTick at 0x7ce0dc7eace0>,
<matplotlib.axis.XTick at 0x7ce0dc7fbcd0>,
<matplotlib.axis.XTick at 0x7ce0dc80cb20>],
[Text(0, 0, '0'),
Text(5, 0, '5'),
Text(10, 0, '10'),
Text(15, 0, '15'),
Text(20, 0, '20')])
```



```
# We create the knn and we train and obtain results
n_neighbors = 7

knn = KNeighborsClassifier(n_neighbors)
knn.fit(X_train, y_train)
print('Accuracy of K-NN classifier on training set: {:.2f}'
      .format(knn.score(X_train, y_train)))
print('Accuracy of K-NN classifier on test set: {:.2f}'
      .format(knn.score(X_test, y_test)))
```

```
Accuracy of K-NN classifier on training set: 0.90
Accuracy of K-NN classifier on test set: 0.86
```

```
pred = knn.predict(X_test)
print(confusion_matrix(y_test, pred))
print(classification_report(y_test, pred))
```

```
[ 9  0  1  0  0]
[ 0  1  0  0  0]
[ 0  1 17  0  1]
[ 0  0  2  8  0]
[ 0  0  4  0 21]]
```

	precision	recall	f1-score	support
1	1.00	0.90	0.95	10
2	0.50	1.00	0.67	1
3	0.71	0.89	0.79	19
4	1.00	0.80	0.89	10
5	0.95	0.84	0.89	25
accuracy			0.86	65
macro avg	0.83	0.89	0.84	65
weighted avg	0.89	0.86	0.87	65

```
#Here we can check the results
h = .02 # step size in the mesh

# Create color maps
```

```
cmap_light = ListedColormap(['#FFAAAA', '#ffcc99', '#ffffb3', '#b3ffff', '#c2f0c2'])
cmap_bold = ListedColormap(['#FF0000', '#ff9933', '#FFFF00', '#00ffff', '#00FF00'])
```

```
# we create an instance of Neighbours Classifier and fit the data.
clf = KNeighborsClassifier(n_neighbors, weights='distance')
clf.fit(X, y)
```

```
# Plot the decision boundary. For that, we will assign a color to each
# point in the mesh [x_min, x_max]x[y_min, y_max].
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])
```

```
# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.figure()
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
```

```
# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold,
            edgecolor='k', s=20)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
```

```
patch0 = mpatches.Patch(color='#FF0000', label='1')
patch1 = mpatches.Patch(color='#ff9933', label='2')
patch2 = mpatches.Patch(color='#FFFF00', label='3')
patch3 = mpatches.Patch(color='#00ffff', label='4')
patch4 = mpatches.Patch(color='#00FF00', label='5')
plt.legend(handles=[patch0, patch1, patch2, patch3, patch4])
```

```
plt.title("5-Class classification (k = %i)"
          % (n_neighbors))
```

```
plt.show()
```

## 5-Class classification (k = 7)

A k-NN (k-Nearest Neighbors) classifier from the scikit-learn library does not use a loss function or an optimizer in the way that, for example, a neural network would. k-NN classifies points according to the majority of votes of the k nearest neighbors, and there is no iterative optimization process involved as in other learning models.

If you want to add a loss function and an optimizer, you would need to switch to a different model that requires them, such as a neural network.

